

A Comprehensive Guide to Testing FastAPI and SQLAlchemy 2.0 with pytest-asyncio

This report provides a comprehensive analysis of architectural and implementation strategies for testing applications built with FastAPI and SQLAlchemy 2.0 in an asynchronous Python environment. Addressing the core user requirement for a robust, production-parity testing solution that facilitates architectural changes, this guide synthesizes best practices from extensive research into **pytest-asyncio**, async database interactions, and transactional isolation. It moves beyond simple solutions to offer a deep dive into managing event loops, session lifecycles, and test parallelization, culminating in a recommended architecture that balances performance, reliability, and code maintainability.

Foundational Architecture for Async Testing with pytest-asyncio

Establishing a reliable testing foundation is paramount when dealing with asynchronous code, as improper setup can lead to non-deterministic failures, resource leaks, and misleading test results. The core challenge lies in managing the lifecycle of the asyncio event loop and ensuring all asynchronous resources are correctly initialized and torn down. The most effective approach involves a deliberate configuration of the testing framework and the use of fixtures to orchestrate the application's state. The primary tools for this task are **pytest** itself, extended with **pytest-asyncio** to enable support for **async def** test functions^{1 55}. This plugin integrates deeply with pytest's fixture system to provide scoped event loops, which are essential for running asynchronous code within synchronous test runners¹⁷.

A critical decision point in this foundational architecture is the configuration of the event loop. While it may be tempting to create a single, long-lived session-scoped event loop to maximize performance, this practice is fraught with peril. Session-scoped event loops are prone to causing "RuntimeError: Event loop is closed" or "Future attached to a different loop" errors, especially when combined with plugins like **pytest-xdist** for parallel test execution^{12 43}. These errors arise because each worker process or thread spawned by **xdist** gets its own event loop, leading to conflicts when a resource created in one loop is used in another²⁸. Furthermore, even without parallelization, holding a strong reference to a task in a session-scoped loop prevents garbage collection, potentially leading to unbounded memory growth^{26 32}. Therefore, the recommended pattern is to avoid creating a custom, session-scoped **event_loop** fixture altogether. The **pytest-asyncio** plugin manages this automatically and will create scoped loops (function, class, module) as needed for tests and fixtures⁴⁹.

Instead of a session-scoped loop, the focus should be on using scoped fixtures for the application's key dependencies. For example, a database engine fixture should typically have a function scope to ensure a clean slate for each test, while other less expensive resources might be shared at the module

level ^{45 53}. When using **pytest-asyncio** version 1.0.0 and later, the plugin enforces stricter rules and removes the deprecated **event_loop** fixture, forcing users to adopt the new, more explicit loop management model based on **loop_scope** markers ⁴⁹. This change, while requiring migration effort, leads to more predictable and reliable test runs. To summarize, the foundational architecture should prioritize safety and isolation over premature optimization. The strategy involves relying on **pytest-asyncio** to manage event loops, avoiding custom session-scoped event loop fixtures, and carefully selecting fixture scopes to balance performance and test independence. This disciplined approach prevents common pitfalls and creates a stable platform upon which more complex testing patterns, such as transactional isolation, can be reliably built.

Implementing Transactional Test Isolation with Async SQLAlchemy

Transactional test isolation is a cornerstone of modern integration testing, providing a fast and reliable way to ensure each test runs against a clean database state without the overhead of recreating schemas or populating data from scratch. In an asynchronous context with SQLAlchemy 2.0, implementing this pattern requires leveraging specific features designed for concurrent operations. The goal is to begin a transaction before a test runs, allow the test to make arbitrary changes to the database, and then roll back those changes after the test completes, effectively leaving no trace ^{3 37}.

The standard and most performant method for achieving this with SQLAlchemy 2.0 is to use nested transactions via savepoints ^{9 37}. This technique is superior to truncating or dropping tables between tests, as it is significantly faster and preserves table auto-increment counters ⁵. The implementation hinges on the **join_transaction_mode="create_savepoint"** argument when yielding the **AsyncSession** from a fixture ¹¹. This tells SQLAlchemy to participate in the existing database transaction (started by the outer fixture) but to create a savepoint specifically for its own operations. If the session flushes any pending objects, these changes are only visible within this nested transaction. Once the test function completes and the session is closed, the outer transaction is rolled back, which in turn rolls back the savepoint and undoes all changes made by the test. This entire process happens on the database server, making it highly efficient. A concrete example of this pattern can be found in a gist demonstrating a working solution where a **session** fixture uses this mode to ensure automatic rollback after each test, even if **await session.commit()** is called within the application code under test ¹¹.

For databases that do not support savepoints, such as SQLite when using the **aiosqlite** driver, an alternative approach is to manually execute **DELETE FROM table_name** statements on all relevant tables after each test ⁵. This method is less elegant but serves the same purpose of cleaning the database state. A more advanced pattern for PostgreSQL involves using **TRUNCATE CASCADE** to remove all data from tables, which is also very fast and respects foreign key constraints ¹⁹. However, care must be taken to execute these commands in the reverse order of table dependencies to avoid foreign key constraint violations ⁵. Regardless of the method chosen, the key principle remains the same: encapsulate the database connection and transaction management within a pytest fixture. This fixture then becomes the provider of the database session dependency for the tests,

ensuring that every interaction goes through this isolated channel. By adopting this pattern, developers can write fast, reliable integration tests that verify the behavior of their domain logic and database interactions without polluting the database or suffering from slow test execution times.

Isolation Method	Description	Performance	Database Support	Key Considerations
Nested Transactions (Savepoints)	Uses SAVEPOINT functionality of the database. The session joins the outer transaction but creates its own savepoint. The outer transaction is rolled back after the test. ^{9 11}	High - Very fast as it's managed server-side.	Most major RDBMS (PostgreSQL, MySQL). Not supported by SQLite/aiosqlite. ³³	Requires <code>join_transaction_mode="create_savepoint"</code> in the session fixture. Ensures <code>session.commit()</code> is also rolled back.
Table Truncation / Deletion	Manually deletes all rows from tested tables (DELETE FROM table_name) or uses TRUNCATE CASCADE . ^{5 19}	Medium to High - Faster than recreating schema. TRUNCATE is very fast.	All databases, including SQLite. ⁵	Must be careful of table dependencies to avoid Foreign Key errors. Preserves table structure and <code>AUTO_INCREMENT</code> values.
Schema Recreation	Drops and recreates the entire database schema before each test. ^{20 41}	Low - Slowest option due to DDL execution.	All databases. ⁴¹	Guarantees absolute isolation but is the most time-consuming method. Counterintuitively, in some async setups, it can be found to be faster than rollbacks ⁴¹ .

Managing Dependencies and Application State in Async Tests

In a well-architected FastAPI application, business logic resides in service or repository layers that are independent of the web framework. This separation is crucial for writing clear and maintainable unit and integration tests. When testing with **pytest**, the primary mechanism for injecting dependencies and mocking external services is through dependency overrides. This allows the test suite to control

the exact instances of classes or functions that the code under test interacts with, ensuring predictable behavior and isolating the test from external systems ^{6 15}.

The most direct way to achieve this in a FastAPI application is by using the **app.dependency_overrides** dictionary. In a pytest fixture that provides an **AsyncClient** for making requests to the application, this dictionary can be populated to replace the standard dependency injection graph with test-specific implementations ^{15 20}. For example, instead of fetching a database session from a real **get_db** dependency, the override can provide a fixture-scoped list that acts as an in-memory database. This approach is particularly effective for unit tests where the goal is to verify the logic of a single function or class in isolation. For instance, a test for an API endpoint could mock the underlying data access layer (DAL) or repository to return predefined data, allowing the test to focus solely on the endpoint's request handling and response formatting logic ¹⁰.

When testing integration scenarios that involve the full stack, including the database, dependency overrides are still used, but in a different manner. Here, the override would point to a fixture that provides a database session within a transactional isolation block ^{3 11}. This allows the test to make requests that interact with the real database, but the changes are safely rolled back afterward. An alternative to dependency overrides is to leverage environment variables to configure the application differently for the test environment ¹⁹. For example, setting an environment variable like **TEST=1** could cause the application's startup event to connect to a different database URL, thus preventing tests from interfering with the development or production databases ³. Both approaches are valid; dependency overrides are more granular and flexible for swapping out individual components, while environment-based separation is simpler for large-scale configuration changes.

Mocking is another essential tool for managing dependencies, especially for external services or I/O-bound operations. For asynchronous functions, Python's built-in **unittest.mock.AsyncMock** should be used to create mock objects that can be awaited ^{9 54}. Libraries like **aiomock** also provide an **AsyncMock** class ¹⁰. A practical example demonstrates mocking an async method of a **CatFact** class to return a predetermined value, thereby avoiding a slow and unreliable real HTTP call to an external API ¹⁰. This ensures that the test's execution time and success are not dependent on the availability or performance of an external service. By combining dependency injection with strategic mocking and transactional database isolation, a powerful and versatile testing ecosystem can be created. This ecosystem supports a wide range of testing styles, from pure unit tests of isolated logic to end-to-end integration tests that validate the entire application workflow, all while maintaining control over application state and external dependencies.

Navigating Common Pitfalls: Event Loop Management and Memory Leaks

While the foundational architecture and testing patterns provide a solid base, navigating the common pitfalls of async testing requires vigilance and a deep understanding of how the asyncio event loop, tasks, and resources are managed. Two of the most prevalent issues are event loop-related runtime errors and memory leaks stemming from improperly managed background tasks or references.

Event loop errors are often the most immediate and frustrating problems encountered. Errors like "RuntimeError: Event loop is closed," "RuntimeError: Task attached to a different loop," or "RuntimeError: no running event loop" are symptoms of mismatched fixture scopes or incorrect usage of async clients^{12 43}. As established, the root cause is frequently a session-scoped fixture (like a database engine or an HTTP client) being created in one event loop and then used in tests that run in a different loop, especially under **pytest-xdist**^{28 44}. The definitive solution is to avoid session-scoped fixtures for objects tied to a specific event loop and to use **@pytest.mark.asyncio(loop_scope="session")** on tests or fixtures that must share a loop, though this is generally discouraged^{43 51}. A more robust pattern is to use **httpx.AsyncClient** with **ASGITransport** as a session-scoped fixture, which properly manages the ASGI application's lifecycle and ensures consistent event loop usage^{42 44}. Another subtle issue is the creation of long-running tasks that are not properly tracked. Holding a reference to a **asyncio.Task** object can prevent the garbage collector from reclaiming it, leading to unbounded memory growth²⁶. The correct pattern is to schedule a task and discard the reference immediately, optionally tracking it in a global set with a done callback to keep it alive only until completion²⁶.

Memory leaks in async Python applications, which manifest during testing just as they do in production, can be more insidious. They often stem from objects retaining references to exceptions or other data structures that include tracebacks, which themselves hold references to local variables³¹. One common leak occurs with **aiohttp.ClientSession** if it is not explicitly closed, as the connector can hold open sockets and event loops³¹. Using **aiohttp.TCPConnector(force_close=True)** can help mitigate this. Another source is the use of **run_in_executor**, which can lead to memory bloat if futures are not awaited promptly, as the underlying thread pool executor does not shut down automatically³⁵. Profiling tools like **tracemalloc** can be invaluable for identifying the source of these leaks³¹. The provided context highlights several other potential leak sources, including SSL contexts in **aiohttp**, failed cleanup of weak references, and circular references involving partial functions³¹. For testing, a proactive approach involves instrumenting the test runner to monitor memory usage and using tools to detect reference cycles. By understanding these common failure modes and adopting defensive coding and testing practices—such as always awaiting tasks, properly closing async resources, and monitoring memory—it is possible to build a resilient testing suite that does not introduce instability into the development process.

Advanced Strategies for Scalability and Production Parity

As an application and its test suite grow, achieving a balance between fast feedback cycles and high-fidelity testing becomes increasingly challenging. The default patterns, while robust, can become bottlenecks. Advanced strategies address these scalability concerns while striving to maintain production parity, ensuring that tests remain a true reflection of the application's behavior.

One of the most significant performance drains in async testing is the overhead of repeatedly creating and tearing down database connections and sessions. A powerful optimization is to pre-warm the connection pool in the async engine fixture²⁵. By establishing a number of connections

equal to the pool size (e.g., 10) at the start of the test session, subsequent tests can reuse these idle connections rather than paying the cost of creating new ones for each session-scoped fixture or test²⁵. This dramatically improves throughput in high-concurrency scenarios. Similarly, for applications that heavily rely on third-party APIs, using a library like **aioresponses** or **aresponses** can provide highly performant and deterministic mocking of HTTP calls, shielding the tests from network latency and external service failures⁷.

For achieving ultimate production parity, especially in distributed environments, Testcontainers are an excellent choice. Testcontainers allow you to spin up disposable, identical instances of services like PostgreSQL, Redis, or Kafka within Docker containers as part of your test suite³⁸. This ensures that the application is tested against a real database backend with the same version and configuration as production, including its specific async driver (e.g., **asyncpg**). This approach eliminates the risk of tests passing locally with one driver (like SQLite's async adapter) and failing in CI with another. The trade-off is increased complexity and slower test execution compared to in-memory databases, but the assurance of production-like behavior is often worth the investment⁵¹³.

Parallel test execution with **pytest-xdist** is another key strategy for scaling test suites⁵⁰. By distributing tests across multiple CPU cores, total execution time can be significantly reduced. However, parallelism introduces its own challenges, primarily around shared state. As discussed previously, session-scoped resources can cause conflicts. The most robust solution for parallelizing tests that interact with a database is to have each worker process create its own isolated test database. This can be achieved by modifying the test database URL to include a unique identifier (e.g., a UUID) for each worker, ensuring complete isolation⁵. Finally, for applications with complex, concurrent database behaviors, it is beneficial to write tests that specifically exercise these scenarios. PostgreSQL's MVCC and snapshot isolation levels can lead to phenomena like non-repeatable reads and phantom reads, which can cause logical inconsistencies if not handled correctly²². Writing tests that simulate conflicting concurrent transactions (e.g., using **asyncio.gather**) against different isolation levels (**REPEATABLE READ**, **SERIALIZABLE**) can uncover subtle bugs that are otherwise hard to detect²³. These advanced strategies transform the testing suite from a passive verifier into an active tool for improving the application's resilience and performance under realistic conditions.

Recommended Architecture and Best Practices

Synthesizing the preceding analysis, a recommended architecture for testing a FastAPI and SQLAlchemy 2.0 application prioritizes production parity, test isolation, and long-term maintainability. This architecture is built on a solid foundation of **pytest** and **pytest-asyncio**, augmented with best-practice patterns for database interaction and dependency management.

The core of this architecture rests on two pillars: a robust event loop management strategy and a reliable transactional isolation pattern. First, the event loop should never be manually managed with a session-scoped fixture. Instead, rely on **pytest-asyncio**'s automatic, scoped loop creation and adopt the **auto** mode in **pytest.ini** for simplicity¹⁷. This minimizes the risk of common

runtime errors. Second, for database testing, the preferred pattern is transactional isolation using nested transactions via savepoints^{9 11}. This is achieved by creating a function-scoped `session` fixture that uses `join_transaction_mode="create_savepoint"`¹¹. This approach is fast, guarantees a clean state for every test, and works seamlessly with `session.commit()` calls in the application code.

Dependency management should favor dependency injection and mocking over monolithic test setups. Use FastAPI's `Depends` mechanism liberally and leverage `app.dependency_overrides` in tests to inject mock services, in-memory repositories, or transactional database sessions^{6 15}. This keeps tests focused and decoupled. For external services, prefer mocking with `AsyncMock` to avoid flaky tests dependent on network conditions¹⁰. This aligns with the principle of isolating the unit of work.

To conclude, the following table summarizes the key architectural decisions and their rationale:

Component	Recommended Pattern	Rationale	Supporting Context
Event Loop	Let <code>pytest-asyncio</code> manage scoped loops automatically.	Avoids common "loop closed" errors and memory leaks associated with manual management.	43 49
Database Engine	Function-scoped fixture with <code>NullPool</code> for tests.	Prevents connection pooling issues in tests and ensures a clean state per test.	40
Session Isolation	Function-scoped <code>session</code> fixture with <code>join_transaction_mode="create_savepoint"</code> .	Provides fast, reliable, and production-like isolation.	9 11
HTTP Client	Session-scoped <code>AsyncClient</code> with <code>ASGITransport</code> .	Ensures consistent event loop usage and proper ASGI app lifecycle management.	42 44
Dependency Override	Use <code>app.dependency_overrides</code> for mocking services/repositories.	Decouples tests from concrete	15 20

Component	Recommended Pattern	Rationale	Supporting Context
		implementations, enabling focused unit and integration tests.	
Test Parallelization	Use pytest-xdist with isolated test databases per worker.	Maximizes test execution speed while preserving test integrity and isolation.	5 50

By adhering to this architecture and its underlying principles, teams can build a testing suite that not only verifies correctness but also serves as a safe playground for refactoring and evolving the application's architecture. This approach directly addresses the user's need for an architectural solution that simplifies testing by isolating commits and ensures that the application is rigorously tested under conditions that mirror its production environment.

Reference

1. Mastering pytest-asyncio and Event Loops for FastAPI and ... <https://medium.com/@connect.hashblock/async-testing-with-pytest-mastering-pytest-asyncio-and-event-loops-for-fastapi-and-beyond-37c613f1cfa3>
2. Transactional Unit Tests with Pytest and Async SQLAlchemy <https://www.core27.co/post/transactional-unit-tests-with-pytest-and-async-sqlalchemy>
3. Setting up a FastAPI App with Async SQLALchemy 2.0 & ... <https://medium.com/@tclaitken/setting-up-a-fastapi-app-with-async-sqlalchemy-2-0-pydantic-v2-e6c540be4308>
4. Asynchronous I/O (asyncio) — SQLAlchemy 2.0 ... <http://docs.sqlalchemy.org/en/latest/orm/extensions/asyncio.html>
5. Fastapi, async SQLAlchemy, pytest, and Alembic (all using ... <https://thedmitry.pw/blog/2023/08/fastapi-async-sqlalchemy-pytest-and-alembic/>
6. Build an async python service with FastAPI & SQLAlchemy <https://towardsdatascience.com/build-an-async-python-service-with-fastapi-sqlalchemy-196d8792fa08/>
7. [timofurrer/awesome-asyncio](https://github.com/timofurrer/awesome-asyncio): A curated list of ... <https://github.com/timofurrer/awesome-asyncio>
8. A Battle of Async Titans: Django ORM Async vs. SQLAlchemy <https://distillery.com/blog/a-battle-of-async-titans-django-orm-async-vs-sqlalchemy-async/>

9. Flask-SQLAlchemy and Async Programming <https://moldstud.com/articles/p-flask-sqlalchemy-and-async-programming-essential-insights-you-need-to-know>
10. A Practical Guide To Async Testing With Pytest-Asyncio <https://pytest-with-eric.com/pytest-advanced/pytest-asyncio/>
11. Pytest + FastAPI + Async SQLAlchemy <https://gist.github.com/e-kondr01/969ae24f2e2f31bd52a81fa5a1fe0f96>
12. Creating Pytest tests for async API calls with FastAPI and ... <https://stackoverflow.com/questions/78612429/creating-pytest-tests-for-async-api-calls-with-fastapi-and-mongodb>
13. Developing and Testing an Asynchronous API with FastAPI ... <https://testdriven.io/blog/fastapi-crud/>
14. FastAPI and async SQLAlchemy 2.0 with pytest done right <https://praciano.com.br/fastapi-and-async-sqlalchemy-20-with-pytest-done-right.html>
15. Testing - JetBrains Guide <https://www.jetbrains.com/guide/python/tutorials/fastapi-aws-kubernetes/testing/>
16. pytest-asyncio has a closed event loop, but only when ... <https://stackoverflow.com/questions/61022713/pytest-asyncio-has-a-closed-event-loop-but-only-when-running-all-tests>
17. Concepts — pytest-asyncio 1.2.0 documentation <https://pytest-asyncio.readthedocs.io/en/stable/concepts.html>
18. Session Basics — SQLAlchemy 2.0 Documentation http://docs.sqlalchemy.org/en/latest/orm/session_basics.html
19. Testing FastAPI with async database session <https://dev.to/whchi/testing-fastapi-with-async-database-session-1b5d>
20. Patterns and Practices for using SQLAlchemy 2.0 with FastAPI <https://chaoticengineer.hashnode.dev/fastapi-sqlalchemy>
21. FastAPI with Async SQLAlchemy, SQLAlchemy, and Alembic <https://testdriven.io/blog/fastapi-sqlmodel/>
22. Transaction Isolation in Postgres, explained <https://www.thenile.dev/blog/transaction-isolation-postgres>
23. Does SQLAlchemy start DB transaction in session? <https://stackoverflow.com/questions/76414845/does-sqlalchemy-start-db-transaction-in-session>
24. Essential pytest asyncio Tips for Modern Async Testing <https://blog.mergify.com/pytest-asyncio-2/>
25. Performance issues with async postgresql #8137 <https://github.com/sqlalchemy/sqlalchemy/discussions/8137>
26. Asyncio with memory leak (Python) <https://stackoverflow.com/questions/69521780/asyncio-with-memory-leak-python>

27. Spring Boot Transaction Management, Propagation, Isolation ... <https://beratyesbek.medium.com/spring-boot-transaction-management-propagation-isolation-levels-f3981b22ef4d>
28. pytest issues with a session scoped fixture and asyncio <https://stackoverflow.com/questions/63713575/pytest-issues-with-a-session-scoped-fixture-and-asyncio>
29. Asynchronous Database Sessions in FastAPI with ... <https://dev.to/akarshan/asynchronous-database-sessions-in-fastapi-with-sqlalchemy-1o7e>
30. Async fixtures may break current event loop · Issue #868 <https://github.com/pytest-dev/pytest-asyncio/issues/868>
31. 7 AsyncIO Memory Leaks Silently Destroying Production ... <https://python.plainenglish.io/7-asyncio-memory-leaks-silently-destroying-production-python-systems-fix-before-its-too-late-4cf724ea1174>
32. Why does asyncio.wait keep a task with a reference around ... <https://stackoverflow.com/questions/73975798/why-does-asyncio-wait-keep-a-task-with-a-reference-around-despite-exceeding-the>
33. SQLAlchemy 2.0 - GINO 1.1.0b2 documentation <https://python-gino.org/docs/en/1.1b2/explanation/sa20.html>
34. "pytest-asyncio will close the event loop for you, but future ... <https://github.com/pytest-dev/pytest-asyncio/issues/531>
35. Potential memory leak with asyncio and run_in_executor <https://bugs.python.org/issue41699>
36. SQL (Relational) Databases <https://fastapi.tiangolo.com/tutorial/sql-databases/>
37. How to set up and tear down a database between tests in ... <https://stackoverflow.com/questions/67255653/how-to-set-up-and-tear-down-a-database-between-tests-in-fastapi>
38. FastAPI Best Practices and Conventions we used at our ... <https://github.com/zhanyamkanov/fastapi-best-practices>
39. Mastering SQLAlchemy: A Comprehensive Guide for ... <https://medium.com/@ramanbazhanau/mastering-sqlalchemy-a-comprehensive-guide-for-python-developers-ddb3d9f2e829>
40. Fast and furious: async testing with FastAPI and pytest <https://weirdsheeplabs.com/blog/fast-and-furious-async-testing-with-fastapi-and-pytest>
41. SQLAlchemy + pytest + asyncio optimization <https://blacksheephacks.pl/sqlalchemy-pytest-asyncio-optimization/>
42. How to set event loop while testing app using pytest-asyncio <https://github.com/fastapi/fastapi/discussions/8415>
43. FastAPI & Pytest. Got Future attached to a different loop <https://stackoverflow.com/questions/79158433/fastapi-pytest-got-future-attached-to-a-different-loop>

44. Event loop is closed" when using pytest-asyncio to test ... <https://stackoverflow.com/a/73811894/16347180>
45. Boost Test Speed with Pytest Fixture Scope <https://blog.mergify.com/pytest-fixture-scope/>
46. Python's asyncio: A Hands-On Walkthrough <https://realpython.com/async-io-python/>
47. What Are Pytest Fixture Scopes? (How To Choose The ... <https://pytest-with-eric.com/fixtures/pytest-fixture-scope/>
48. pytest fixtures: explicit, modular, scalable <https://docs.pytest.org/en/6.2.x/fixture.html>
49. A Deep Dive into pytest-asyncio 1.0 and Migration Strategies <https://thinhdanggroup.github.io/pytest-asyncio-v1-migrate/>
50. Pytest Mastery: Advanced Testing Techniques in Python <https://python.plainenglish.io/pytest-mastery-unleashing-advanced-testing-techniques-in-python-78622036f3ab>
51. Session scoped event loop not actually session scope #944 <https://github.com/pytest-dev/pytest-asyncio/issues/944>
52. Asynchronous Programming with FastAPI: Building ... <https://dev.to/dhruमितdk/asynchronous-programming-with-fastapi-building-efficient-apis-nj1>
53. What Are Pytest Fixture Scopes? (How To Choose The Best ... [https://dag7.it/appunti/dev/Pytest/What-Are-Pytest-Fixture-Scopes-\(How-To-Choose-The-Best-Scope-For-Your-Test\)](https://dag7.it/appunti/dev/Pytest/What-Are-Pytest-Fixture-Scopes-(How-To-Choose-The-Best-Scope-For-Your-Test))
54. async test patterns for Pytest - Anthony Shaw <https://tonybaloney.github.io/posts/async-test-patterns-for-pytest-and-unittest.html>
55. Automating tests for async services using pytest <https://wearecommunity.io/communities/india-devtestsecops-community/articles/1745>